Title: On Real-time Decoding Latency Using Magic State Distillation

Year: 2024

Authors: J. Banerjee, P. Kowalski, P. Dubois, B. Chen (Institute for Advanced Computing Research)

Keywords: Hamiltonian simulation, quantum annealing, quantum circuits, quantum gates, entanglement

Abstract:

Real-time decoding latency remains a critical challenge for quantum error correction, particularly in molecular ground-state estimation. This paper proposes a framework built on magic state distillation that explicitly accounts for circuit depth minimization. We evaluate the framework on simulated Hamiltonians of H2 and LiH molecules and compare it against 5 baseline methods. Experimental results show that the proposed approach improves logical error rate by 19.0% while keeping gate fidelity within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

The rapid maturation of quantum error correction has changed how researchers and practitioners approach real-time decoding latency, yet practical deployment is still limited by scalability, cost, and reproducibility concerns. Over the past decade, magic state distillation have attracted sustained attention as a promising route toward solving real-time decoding latency in realistic settings such as molecular ground-state estimation. Despite considerable progress, existing solutions for real-time decoding latency tend to degrade when moving from controlled laboratory conditions to production workloads, which motivates a more systematic investigation. A central difficulty is that improvements in approximation ratio frequently come at the expense of circuit success probability, forcing designers to accept trade-offs that are poorly understood.

This paper addresses these gaps by studying how magic state distillation interact with quantum phase estimation routines when applied to real-time decoding latency. Our work is motivated by the growing demand for molecular ground-state estimation, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of real-time decoding latency, a practical design built on magic state distillation, and an empirical evaluation on simulated Hamiltonians of H2 and LiH molecules. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Nakamura et al. (2017) introduced an early framework based on magic state distillation, reporting a 15.2% improvement in gate fidelity over previous baselines. A survey by Mukherjee et al. (2012) categorised existing approaches to real-time decoding latency into three families and highlighted the lack of standard benchmarks. Researchers working on quantum error correction have explored quantum phase estimation routines as an alternative, with Takahashi et al. (2010) noting that the approach scales poorly beyond 32 components. Kowalski et al. (2015) analysed circuit depth minimization and showed that purely heuristic solutions cannot guarantee consistent approximation ratio. More recently, OConnor et al. (2013) combined magic state distillation with quantum phase estimation routines, achieving a 24.1% reduction in overhead on simulated Hamiltonians of H2 and LiH molecules.

Several studies have criticised the reliance on synthetic workloads, arguing that results on simulated Hamiltonians of H2 and LiH molecules may not generalise to molecular ground-state estimation. In contrast to these efforts, our work treats real-time decoding latency and circuit depth minimization jointly rather than as independent concerns. Industry reports corroborate these academic findings, pointing to real-time decoding

latency as a persistent obstacle for molecular ground-state estimation. Despite this body of literature, no prior study has systematically compared magic state distillation and quantum phase estimation routines under identical experimental conditions.

3. Methodology

We formalise real-time decoding latency as a constrained optimisation task in which the objective is to maximise logical error rate while keeping approximation ratio within a prescribed budget. The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on magic state distillation, and a validation layer that incorporates quantum phase estimation routines. All experiments were conducted on simulated Hamiltonians of H2 and LiH molecules, which provides 500 labelled instances or configurations covering a broad range of operating conditions. To ensure reproducibility, every experiment was repeated 20 times with independent random seeds, and we report mean values together with 95% confidence intervals.

Hyperparameters were selected through grid search over a validation split, and the final configuration was fixed before any evaluation on held-out data. The core routine applies magic state distillation to reduce the effect of circuit depth minimization, followed by a calibration step that stabilises logical error rate. Our implementation is written in Python and C++, and was executed on a workstation with 64 GB of memory and a single mid-range accelerator. We compare against 5 baseline methods, including a conventional approach built on quantum phase estimation routines and two recent variants reported in the literature.

The proposed design improves logical error rate by 19.0% relative to the strongest baseline and reduces gate fidelity by 32.3%, a statistically significant difference ($p < 0.05$). A sensitivity analysis over 500 parameter settings shows that performance remains stable, with variation in circuit success probability below 11.1%. Threats to validity include the limited diversity of simulated Hamiltonians of H2 and LiH molecules and the possibility that deployment in molecular ground-state estimation introduces conditions not captured in our setup.

4. Conclusion

This paper investigated real-time decoding latency through the lens of magic state distillation, and demonstrated that a carefully designed pipeline can improve logical error rate by 19.0% on simulated Hamiltonians of H2 and LiH molecules. The findings suggest that quantum error correction will continue to benefit from closer integration of magic state distillation and quantum phase estimation routines, particularly in molecular ground-state estimation. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of circuit depth minimization. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for real-time decoding latency.