Title: Cryptographic Hash Functions Meets Lattice-based Encryption Schemes: A Framework for Privacy-preserving Data Sharing

Year: 2021

Authors: M. Nakamura, Y. OConnor, I. Novak, T. Singh (Centre for Secure Systems, Eastvale Institute of Technology)

Keywords: privacy, provable security, zero-knowledge proofs, authentication, digital signatures

Abstract:

Privacy-preserving data sharing remains a critical challenge for cryptographic hash functions, particularly in privacy-preserving machine learning. This paper proposes a framework built on lattice-based encryption schemes that explicitly accounts for ciphertext expansion reduction. We evaluate the framework on a corpus of TLS handshake traces and compare it against 3 baseline methods. Experimental results show that the proposed approach improves key generation latency by 18.9% while keeping ciphertext overhead within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

The rapid maturation of cryptographic hash functions has changed how researchers and practitioners approach privacy-preserving data sharing, yet practical deployment is still limited by scalability, cost, and reproducibility concerns. Over the past decade, lattice-based encryption schemes have attracted sustained attention as a promising route toward solving privacy-preserving data sharing in realistic settings such as privacy-preserving machine learning. A central difficulty is that improvements in encryption throughput frequently come at the expense of security level in bits, forcing designers to accept trade-offs that are poorly understood. This paper addresses these gaps by studying how lattice-based encryption schemes interact with side-channel trace analysis when applied to privacy-preserving data sharing.

Our work is motivated by the growing demand for privacy-preserving machine learning, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of privacy-preserving data sharing, a practical design built on lattice-based encryption schemes, and an empirical evaluation on a corpus of TLS handshake traces. Understanding ciphertext expansion reduction is equally important, because it shapes both the theoretical limits and the operational cost of any proposed approach. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Sharma et al. (2013) introduced an early framework based on lattice-based encryption schemes, reporting a 31.3% improvement in encryption throughput over previous baselines. Subsequent work by Iyer et al. (2011) extended this line of research to privacy-preserving machine learning, although the evaluation was restricted to small-scale settings. A survey by Hoffmann et al. (2014) categorised existing approaches to privacy-preserving data sharing into three families and highlighted the lack of standard benchmarks. Researchers working on cryptographic hash functions have explored side-channel trace analysis as an alternative, with Haddad et al. (2014) noting that the approach scales poorly beyond 8 components. Johansson et al. (2010) analysed ciphertext expansion reduction and showed that purely heuristic solutions cannot guarantee consistent security level in bits.

Several studies have criticised the reliance on synthetic workloads, arguing that results on a corpus of TLS handshake traces may not generalise to privacy-preserving machine learning. In contrast to these efforts, our

work treats privacy-preserving data sharing and ciphertext expansion reduction jointly rather than as independent concerns. Industry reports corroborate these academic findings, pointing to privacy-preserving data sharing as a persistent obstacle for privacy-preserving machine learning. Despite this body of literature, no prior study has systematically compared lattice-based encryption schemes and side-channel trace analysis under identical experimental conditions.

3. Methodology

We formalise privacy-preserving data sharing as a constrained optimisation task in which the objective is to maximise key generation latency while keeping key generation latency within a prescribed budget. The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on lattice-based encryption schemes, and a validation layer that incorporates side-channel trace analysis. To ensure reproducibility, every experiment was repeated 10 times with independent random seeds, and we report mean values together with 95% confidence intervals. Hyperparameters were selected through grid search over a validation split, and the final configuration was fixed before any evaluation on held-out data.

The core routine applies lattice-based encryption schemes to reduce the effect of ciphertext expansion reduction, followed by a calibration step that stabilises encryption throughput. Our implementation is written in Python and C++, and was executed on a workstation with 64 GB of memory and a single mid-range accelerator. We compare against 3 baseline methods, including a conventional approach built on side-channel trace analysis and two recent variants reported in the literature. The proposed design improves key generation latency by 18.9% relative to the strongest baseline and reduces security level in bits by 5.6%, a statistically significant difference ($p < 0.05$).

Ablation experiments show that removing the side-channel trace analysis component reduces key generation latency by approximately 7.3%, confirming its contribution. A sensitivity analysis over 1200 parameter settings shows that performance remains stable, with variation in ciphertext overhead below 6.7%. Threats to validity include the limited diversity of a corpus of TLS handshake traces and the possibility that deployment in privacy-preserving machine learning introduces conditions not captured in our setup.

4. Conclusion

This paper investigated privacy-preserving data sharing through the lens of lattice-based encryption schemes, and demonstrated that a carefully designed pipeline can improve key generation latency by 18.9% on a corpus of TLS handshake traces. The findings suggest that cryptographic hash functions will continue to benefit from closer integration of lattice-based encryption schemes and side-channel trace analysis, particularly in privacy-preserving machine learning. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of ciphertext expansion reduction. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for privacy-preserving data sharing.