Title: On Protocol Formal Verification Using Attribute-based Encryption

Year: 2021

Authors: M. Tanaka, E. Hoffmann (Digital Futures Research Group, Harlow University)

Keywords: cryptography, key exchange, side-channel attacks, provable security, digital signatures

Abstract:

Protocol formal verification remains a critical challenge for public-key infrastructure, particularly in privacy-preserving machine learning. This paper proposes a framework built on attribute-based encryption that explicitly accounts for side-channel leakage resistance. We evaluate the framework on power traces from an ARM Cortex-M4 platform and compare it against 3 baseline methods. Experimental results show that the proposed approach improves encryption throughput by 27.2% while keeping key generation latency within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

The rapid maturation of public-key infrastructure has changed how researchers and practitioners approach protocol formal verification, yet practical deployment is still limited by scalability, cost, and reproducibility concerns. Over the past decade, attribute-based encryption have attracted sustained attention as a promising route toward solving protocol formal verification in realistic settings such as privacy-preserving machine learning. A central difficulty is that improvements in ciphertext overhead frequently come at the expense of key generation latency, forcing designers to accept trade-offs that are poorly understood. This paper addresses these gaps by studying how attribute-based encryption interact with secure multiparty computation when applied to protocol formal verification.

Our work is motivated by the growing demand for privacy-preserving machine learning, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of protocol formal verification, a practical design built on attribute-based encryption, and an empirical evaluation on power traces from an ARM Cortex-M4 platform. Understanding side-channel leakage resistance is equally important, because it shapes both the theoretical limits and the operational cost of any proposed approach. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Nakamura et al. (2012) introduced an early framework based on attribute-based encryption, reporting a 31.4% improvement in security level in bits over previous baselines. Subsequent work by Vasquez et al. (2016) extended this line of research to privacy-preserving machine learning, although the evaluation was restricted to small-scale settings. A survey by Hoffmann et al. (2020) categorised existing approaches to protocol formal verification into three families and highlighted the lack of standard benchmarks. Researchers working on public-key infrastructure have explored secure multiparty computation as an alternative, with Vasquez et al. (2017) noting that the approach scales poorly beyond 32 components. Rossi et al. (2008) analysed side-channel leakage resistance and showed that purely heuristic solutions cannot guarantee consistent encryption throughput.

Several studies have criticised the reliance on synthetic workloads, arguing that results on power traces from an ARM Cortex-M4 platform may not generalise to privacy-preserving machine learning. In contrast to these efforts, our work treats protocol formal verification and side-channel leakage resistance jointly rather than as independent concerns. Industry reports corroborate these academic findings, pointing to protocol formal

verification as a persistent obstacle for privacy-preserving machine learning. Despite this body of literature, no prior study has systematically compared attribute-based encryption and secure multiparty computation under identical experimental conditions.

3. Methodology

We formalise protocol formal verification as a constrained optimisation task in which the objective is to maximise encryption throughput while keeping ciphertext overhead within a prescribed budget. The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on attribute-based encryption, and a validation layer that incorporates secure multiparty computation. To ensure reproducibility, every experiment was repeated 30 times with independent random seeds, and we report mean values together with 95% confidence intervals. Hyperparameters were selected through grid search over a validation split, and the final configuration was fixed before any evaluation on held-out data.

The core routine applies attribute-based encryption to reduce the effect of side-channel leakage resistance, followed by a calibration step that stabilises key generation latency. Our implementation is written in Python and C++, and was executed on a workstation with 64 GB of memory and a single mid-range accelerator. We compare against 3 baseline methods, including a conventional approach built on secure multiparty computation and two recent variants reported in the literature. The proposed design improves encryption throughput by 27.2% relative to the strongest baseline and reduces security level in bits by 37.2%, a statistically significant difference ($p < 0.05$).

Ablation experiments show that removing the secure multiparty computation component reduces encryption throughput by approximately 6.0%, confirming its contribution. A sensitivity analysis over 5000 parameter settings shows that performance remains stable, with variation in ciphertext overhead below 7.4%. Threats to validity include the limited diversity of power traces from an ARM Cortex-M4 platform and the possibility that deployment in privacy-preserving machine learning introduces conditions not captured in our setup.

4. Conclusion

This paper investigated protocol formal verification through the lens of attribute-based encryption, and demonstrated that a carefully designed pipeline can improve encryption throughput by 27.2% on power traces from an ARM Cortex-M4 platform. The findings suggest that public-key infrastructure will continue to benefit from closer integration of attribute-based encryption and secure multiparty computation, particularly in privacy-preserving machine learning. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of side-channel leakage resistance. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for protocol formal verification.