Title: Rethinking Quantum-resistant TLS Deployment: A HQC Code-based Key Encapsulation Perspective

Year: 2019

Authors: B. Novak, S. Fernandez (Faculty of Engineering, Meridian Polytechnic)

Keywords: lattice cryptography, Shor's algorithm, quantum key distribution, post-quantum cryptography, NIST PQC

Abstract:

Quantum-resistant TLS deployment remains a critical challenge for the NIST standardization process, particularly in long-term archival encryption. This paper proposes a framework built on HQC code-based key encapsulation that explicitly accounts for harvest-now-decrypt-later risk. We evaluate the framework on the NIST PQC round-three submissions and compare it against 3 baseline methods. Experimental results show that the proposed approach improves decapsulation time by 32.3% while keeping secret key rate within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

The rapid maturation of the NIST standardization process has changed how researchers and practitioners approach quantum-resistant TLS deployment, yet practical deployment is still limited by scalability, cost, and reproducibility concerns. Over the past decade, HQC code-based key encapsulation have attracted sustained attention as a promising route toward solving quantum-resistant TLS deployment in realistic settings such as long-term archival encryption. Despite considerable progress, existing solutions for quantum-resistant TLS deployment tend to degrade when moving from controlled laboratory conditions to production workloads, which motivates a more systematic investigation. This paper addresses these gaps by studying how HQC code-based key encapsulation interact with Shor's factoring algorithm when applied to quantum-resistant TLS deployment.

Our work is motivated by the growing demand for long-term archival encryption, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of quantum-resistant TLS deployment, a practical design built on HQC code-based key encapsulation, and an empirical evaluation on the NIST PQC round-three submissions. Understanding harvest-now-decrypt-later risk is equally important, because it shapes both the theoretical limits and the operational cost of any proposed approach. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Kowalski et al. (2009) introduced an early framework based on HQC code-based key encapsulation, reporting a 43.6% improvement in logical qubit count over previous baselines. Subsequent work by Silva et al. (2008) extended this line of research to long-term archival encryption, although the evaluation was restricted to small-scale settings. A survey by Petrov et al. (2010) categorised existing approaches to quantum-resistant TLS deployment into three families and highlighted the lack of standard benchmarks. Researchers working on the NIST standardization process have explored Shor's factoring algorithm as an alternative, with Nakamura et al. (2015) noting that the approach scales poorly beyond 16 components. Okafor et al. (2013) analysed harvest-now-decrypt-later risk and showed that purely heuristic solutions cannot guarantee consistent decapsulation time.

More recently, Singh et al. (2010) combined HQC code-based key encapsulation with Shor's factoring algorithm, achieving a 8.4% reduction in overhead on the NIST PQC round-three submissions. In contrast to

these efforts, our work treats quantum-resistant TLS deployment and harvest-now-decrypt-later risk jointly rather than as independent concerns. Industry reports corroborate these academic findings, pointing to quantum-resistant TLS deployment as a persistent obstacle for long-term archival encryption. Despite this body of literature, no prior study has systematically compared HQC code-based key encapsulation and Shor's factoring algorithm under identical experimental conditions.

3. Methodology

We formalise quantum-resistant TLS deployment as a constrained optimisation task in which the objective is to maximise decapsulation time while keeping signature size within a prescribed budget. The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on HQC code-based key encapsulation, and a validation layer that incorporates Shor's factoring algorithm. All experiments were conducted on the NIST PQC round-three submissions, which provides 128 labelled instances or configurations covering a broad range of operating conditions. To ensure reproducibility, every experiment was repeated 20 times with independent random seeds, and we report mean values together with 95% confidence intervals.

The core routine applies HQC code-based key encapsulation to reduce the effect of harvest-now-decrypt-later risk, followed by a calibration step that stabilises secret key rate. Our implementation is written in Python and C++, and was executed on a workstation with 64 GB of memory and a single mid-range accelerator. We compare against 3 baseline methods, including a conventional approach built on Shor's factoring algorithm and two recent variants reported in the literature. The proposed design improves decapsulation time by 32.3% relative to the strongest baseline and reduces secret key rate by 38.2%, a statistically significant difference ($p < 0.05$).

Ablation experiments show that removing the Shor's factoring algorithm component reduces decapsulation time by approximately 5.7%, confirming its contribution. A sensitivity analysis over 64 parameter settings shows that performance remains stable, with variation in logical qubit count below 6.8%. Threats to validity include the limited diversity of the NIST PQC round-three submissions and the possibility that deployment in long-term archival encryption introduces conditions not captured in our setup.

4. Conclusion

This paper investigated quantum-resistant TLS deployment through the lens of HQC code-based key encapsulation, and demonstrated that a carefully designed pipeline can improve decapsulation time by 32.3% on the NIST PQC round-three submissions. The findings suggest that the NIST standardization process will continue to benefit from closer integration of HQC code-based key encapsulation and Shor's factoring algorithm, particularly in long-term archival encryption. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of harvest-now-decrypt-later risk. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for quantum-resistant TLS deployment.