Title: On Low-resource Language Adaptation Using Neural Architecture Search

Year: 2015
Authors: F. Takahashi, F. Larsen (School of Information Sciences, Lakeshore University)
Keywords: robustness, deep learning, neural networks, generative models, federated learning

Abstract:

Low-resource language adaptation remains a critical challenge for foundation models, particularly in medical image diagnosis. This paper proposes a framework built on neural architecture search that explicitly accounts for hallucination reduction. We evaluate the framework on MIMIC-IV clinical records and compare it against 6 baseline methods. Experimental results show that the proposed approach improves mean episodic reward by 30.0% while keeping mean episodic reward within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

The rapid maturation of foundation models has changed how researchers and practitioners approach low-resource language adaptation, yet practical deployment is still limited by scalability, cost, and reproducibility concerns. Over the past decade, neural architecture search have attracted sustained attention as a promising route toward solving low-resource language adaptation in realistic settings such as medical image diagnosis. A central difficulty is that improvements in mean episodic reward frequently come at the expense of BLEU score, forcing designers to accept trade-offs that are poorly understood. This paper addresses these gaps by studying how neural architecture search interact with retrieval-augmented generation pipelines when applied to low-resource language adaptation.

Our work is motivated by the growing demand for medical image diagnosis, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of low-resource language adaptation, a practical design built on neural architecture search, and an empirical evaluation on MIMIC-IV clinical records. Understanding hallucination reduction is equally important, because it shapes both the theoretical limits and the operational cost of any proposed approach. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Tanaka et al. (2009) introduced an early framework based on neural architecture search, reporting a 43.4% improvement in F1 score over previous baselines. Subsequent work by Tanaka et al. (2009) extended this line of research to medical image diagnosis, although the evaluation was restricted to small-scale settings. A survey by Nakamura et al. (2012) categorised existing approaches to low-resource language adaptation into three families and highlighted the lack of standard benchmarks. Researchers working on foundation models have explored retrieval-augmented generation pipelines as an alternative, with Tanaka et al. (2014) noting that the approach scales poorly beyond 128 components. OConnor et al. (2012) analysed hallucination reduction and showed that purely heuristic solutions cannot guarantee consistent BLEU score.

More recently, Novak et al. (2009) combined neural architecture search with retrieval-augmented generation pipelines, achieving a 26.0% reduction in overhead on MIMIC-IV clinical records. Several studies have criticised the reliance on synthetic workloads, arguing that results on MIMIC-IV clinical records may not generalise to medical image diagnosis. In contrast to these efforts, our work treats low-resource language adaptation and hallucination reduction jointly rather than as independent concerns. Despite this body of

literature, no prior study has systematically compared neural architecture search and retrieval-augmented generation pipelines under identical experimental conditions.

3. Methodology

We formalise low-resource language adaptation as a constrained optimisation task in which the objective is to maximise mean episodic reward while keeping BLEU score within a prescribed budget. The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on neural architecture search, and a validation layer that incorporates retrieval-augmented generation pipelines. All experiments were conducted on MIMIC-IV clinical records, which provides 32 labelled instances or configurations covering a broad range of operating conditions. To ensure reproducibility, every experiment was repeated 10 times with independent random seeds, and we report mean values together with 95% confidence intervals.

Hyperparameters were selected through grid search over a validation split, and the final configuration was fixed before any evaluation on held-out data. Our implementation is written in Python and C++, and was executed on a workstation with 64 GB of memory and a single mid-range accelerator. We compare against 6 baseline methods, including a conventional approach built on retrieval-augmented generation pipelines and two recent variants reported in the literature. The proposed design improves mean episodic reward by 30.0% relative to the strongest baseline and reduces F1 score by 6.2%, a statistically significant difference ($p < 0.05$).

Ablation experiments show that removing the retrieval-augmented generation pipelines component reduces mean episodic reward by approximately 6.7%, confirming its contribution. A sensitivity analysis over 128 parameter settings shows that performance remains stable, with variation in mean episodic reward below 11.2%. Threats to validity include the limited diversity of MIMIC-IV clinical records and the possibility that deployment in medical image diagnosis introduces conditions not captured in our setup.

4. Conclusion

This paper investigated low-resource language adaptation through the lens of neural architecture search, and demonstrated that a carefully designed pipeline can improve mean episodic reward by 30.0% on MIMIC-IV clinical records. The findings suggest that foundation models will continue to benefit from closer integration of neural architecture search and retrieval-augmented generation pipelines, particularly in medical image diagnosis. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of hallucination reduction. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for low-resource language adaptation.