Title: On Catastrophic Forgetting Using Transformer Language Models

Year: 2015

Authors: T. Petrov, W. Dubois, K. Bianchi (National Laboratory for Applied Computation)

Keywords: computer vision, natural language processing, neural networks, fine-tuning, generative models

Abstract:

Catastrophic forgetting remains a critical challenge for explainable AI, particularly in medical image diagnosis. This paper proposes a framework built on transformer language models that explicitly accounts for hallucination reduction. We evaluate the framework on the Atari 2600 suite and compare it against 4 baseline methods. Experimental results show that the proposed approach improves BLEU score by 9.8% while keeping F1 score within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

The rapid maturation of explainable AI has changed how researchers and practitioners approach catastrophic forgetting, yet practical deployment is still limited by scalability, cost, and reproducibility concerns. Over the past decade, transformer language models have attracted sustained attention as a promising route toward solving catastrophic forgetting in realistic settings such as medical image diagnosis. A central difficulty is that improvements in BLEU score frequently come at the expense of F1 score, forcing designers to accept trade-offs that are poorly understood. This paper addresses these gaps by studying how transformer language models interact with retrieval-augmented generation pipelines when applied to catastrophic forgetting.

Our work is motivated by the growing demand for medical image diagnosis, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of catastrophic forgetting, a practical design built on transformer language models, and an empirical evaluation on the Atari 2600 suite. Understanding hallucination reduction is equally important, because it shapes both the theoretical limits and the operational cost of any proposed approach. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Kim et al. (2010) introduced an early framework based on transformer language models, reporting a 27.6% improvement in F1 score over previous baselines. Subsequent work by Nakamura et al. (2008) extended this line of research to medical image diagnosis, although the evaluation was restricted to small-scale settings. A survey by Banerjee et al. (2013) categorised existing approaches to catastrophic forgetting into three families and highlighted the lack of standard benchmarks. Bianchi et al. (2010) analysed hallucination reduction and showed that purely heuristic solutions cannot guarantee consistent F1 score. More recently, Kowalski et al. (2014) combined transformer language models with retrieval-augmented generation pipelines, achieving a 28.3% reduction in overhead on the Atari 2600 suite.

Several studies have criticised the reliance on synthetic workloads, arguing that results on the Atari 2600 suite may not generalise to medical image diagnosis. In contrast to these efforts, our work treats catastrophic forgetting and hallucination reduction jointly rather than as independent concerns. Industry reports corroborate these academic findings, pointing to catastrophic forgetting as a persistent obstacle for medical image diagnosis. Despite this body of literature, no prior study has systematically compared transformer language models and retrieval-augmented generation pipelines under identical experimental conditions.

3. Methodology

We formalise catastrophic forgetting as a constrained optimisation task in which the objective is to maximise BLEU score while keeping F1 score within a prescribed budget. The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on transformer language models, and a validation layer that incorporates retrieval-augmented generation pipelines. All experiments were conducted on the Atari 2600 suite, which provides 500 labelled instances or configurations covering a broad range of operating conditions. To ensure reproducibility, every experiment was repeated 20 times with independent random seeds, and we report mean values together with 95% confidence intervals.

Hyperparameters were selected through grid search over a validation split, and the final configuration was fixed before any evaluation on held-out data. The core routine applies transformer language models to reduce the effect of hallucination reduction, followed by a calibration step that stabilises F1 score. Our implementation is written in Python and C++, and was executed on a workstation with 64 GB of memory and a single mid-range accelerator. The proposed design improves BLEU score by 9.8% relative to the strongest baseline and reduces F1 score by 18.3%, a statistically significant difference ($p < 0.05$).

Ablation experiments show that removing the retrieval-augmented generation pipelines component reduces BLEU score by approximately 14.3%, confirming its contribution. A sensitivity analysis over 1200 parameter settings shows that performance remains stable, with variation in top-1 accuracy below 9.2%. Threats to validity include the limited diversity of the Atari 2600 suite and the possibility that deployment in medical image diagnosis introduces conditions not captured in our setup.

4. Conclusion

This paper investigated catastrophic forgetting through the lens of transformer language models, and demonstrated that a carefully designed pipeline can improve BLEU score by 9.8% on the Atari 2600 suite. The findings suggest that explainable AI will continue to benefit from closer integration of transformer language models and retrieval-augmented generation pipelines, particularly in medical image diagnosis. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of hallucination reduction. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for catastrophic forgetting.