Title: Zero-knowledge Proof Systems: Toward Post-quantum Migration of Legacy Protocols in Electronic Voting

Year: 2021

Authors: M. Ahmed, T. Petrov, S. Mukherjee, B. Nakamura (Institute for Advanced Computing Research)

Keywords: zero-knowledge proofs, cryptography, homomorphic encryption, privacy, digital signatures

Abstract:

Post-quantum migration of legacy protocols remains a critical challenge for cryptographic hash functions, particularly in electronic voting. This paper proposes a framework built on zero-knowledge proof systems that explicitly accounts for ciphertext expansion reduction. We evaluate the framework on the NIST test vector suites and compare it against 5 baseline methods. Experimental results show that the proposed approach improves key generation latency by 28.8% while keeping security level in bits within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

Over the past decade, zero-knowledge proof systems have attracted sustained attention as a promising route toward solving post-quantum migration of legacy protocols in realistic settings such as electronic voting. Despite considerable progress, existing solutions for post-quantum migration of legacy protocols tend to degrade when moving from controlled laboratory conditions to production workloads, which motivates a more systematic investigation. A central difficulty is that improvements in key generation latency frequently come at the expense of ciphertext overhead, forcing designers to accept trade-offs that are poorly understood. This paper addresses these gaps by studying how zero-knowledge proof systems interact with ML-DSA lattice signatures when applied to post-quantum migration of legacy protocols.

Our work is motivated by the growing demand for electronic voting, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of post-quantum migration of legacy protocols, a practical design built on zero-knowledge proof systems, and an empirical evaluation on the NIST test vector suites. Understanding ciphertext expansion reduction is equally important, because it shapes both the theoretical limits and the operational cost of any proposed approach. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Iyer et al. (2020) introduced an early framework based on zero-knowledge proof systems, reporting a 43.6% improvement in encryption throughput over previous baselines. Subsequent work by Vasquez et al. (2010) extended this line of research to electronic voting, although the evaluation was restricted to small-scale settings. A survey by Silva et al. (2018) categorised existing approaches to post-quantum migration of legacy protocols into three families and highlighted the lack of standard benchmarks. Researchers working on cryptographic hash functions have explored ML-DSA lattice signatures as an alternative, with Kim et al. (2019) noting that the approach scales poorly beyond 8 components. Chen et al. (2018) analysed ciphertext expansion reduction and showed that purely heuristic solutions cannot guarantee consistent encryption throughput.

More recently, Zhang et al. (2013) combined zero-knowledge proof systems with ML-DSA lattice signatures, achieving a 24.6% reduction in overhead on the NIST test vector suites. Several studies have criticised the reliance on synthetic workloads, arguing that results on the NIST test vector suites may not generalise to electronic voting. Industry reports corroborate these academic findings, pointing to post-quantum migration of legacy protocols as a persistent obstacle for electronic voting. Despite this body of literature, no prior study has

systematically compared zero-knowledge proof systems and ML-DSA lattice signatures under identical experimental conditions.

3. Methodology

The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on zero-knowledge proof systems, and a validation layer that incorporates ML-DSA lattice signatures. All experiments were conducted on the NIST test vector suites, which provides 8 labelled instances or configurations covering a broad range of operating conditions. To ensure reproducibility, every experiment was repeated 30 times with independent random seeds, and we report mean values together with 95% confidence intervals. Hyperparameters were selected through grid search over a validation split, and the final configuration was fixed before any evaluation on held-out data.

The core routine applies zero-knowledge proof systems to reduce the effect of ciphertext expansion reduction, followed by a calibration step that stabilises encryption throughput. Our implementation is written in Python and C++, and was executed on a workstation with 64 GB of memory and a single mid-range accelerator. We compare against 5 baseline methods, including a conventional approach built on ML-DSA lattice signatures and two recent variants reported in the literature. The proposed design improves key generation latency by 28.8% relative to the strongest baseline and reduces key generation latency by 33.8%, a statistically significant difference ($p < 0.05$).

Ablation experiments show that removing the ML-DSA lattice signatures component reduces key generation latency by approximately 5.8%, confirming its contribution. A sensitivity analysis over 5000 parameter settings shows that performance remains stable, with variation in ciphertext overhead below 12.9%. Threats to validity include the limited diversity of the NIST test vector suites and the possibility that deployment in electronic voting introduces conditions not captured in our setup.

4. Conclusion

This paper investigated post-quantum migration of legacy protocols through the lens of zero-knowledge proof systems, and demonstrated that a carefully designed pipeline can improve key generation latency by 28.8% on the NIST test vector suites. The findings suggest that cryptographic hash functions will continue to benefit from closer integration of zero-knowledge proof systems and ML-DSA lattice signatures, particularly in electronic voting. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of ciphertext expansion reduction. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for post-quantum migration of legacy protocols.