Title: Hamiltonian Simulation Meets Variational Quantum Eigensolvers: A Framework for Magic State Overhead Reduction

Year: 2023

Authors: P. Alvarez, M. Ahmed (Centre for Secure Systems, Eastvale Institute of Technology)

Keywords: decoherence, entanglement, hybrid algorithms, Hamiltonian simulation, NISQ

Abstract:

Magic state overhead reduction remains a critical challenge for Hamiltonian simulation, particularly in portfolio optimization. This paper proposes a framework built on variational quantum eigensolvers that explicitly accounts for qubit decoherence suppression. We evaluate the framework on a 27-qubit device calibration log and compare it against 6 baseline methods. Experimental results show that the proposed approach improves logical error rate by 13.9% while keeping gate fidelity within acceptable bounds. We further analyse its limitations and discuss implications for practical deployment.

1. Introduction

The rapid maturation of Hamiltonian simulation has changed how researchers and practitioners approach magic state overhead reduction, yet practical deployment is still limited by scalability, cost, and reproducibility concerns. Over the past decade, variational quantum eigensolvers have attracted sustained attention as a promising route toward solving magic state overhead reduction in realistic settings such as portfolio optimization. Despite considerable progress, existing solutions for magic state overhead reduction tend to degrade when moving from controlled laboratory conditions to production workloads, which motivates a more systematic investigation. This paper addresses these gaps by studying how variational quantum eigensolvers interact with surface-code error correction when applied to magic state overhead reduction.

Our work is motivated by the growing demand for portfolio optimization, where reliability and transparency requirements are especially strict. We make three contributions: a unified formulation of magic state overhead reduction, a practical design built on variational quantum eigensolvers, and an empirical evaluation on a 27-qubit device calibration log. Understanding qubit decoherence suppression is equally important, because it shapes both the theoretical limits and the operational cost of any proposed approach. The remainder of this paper reviews related work, describes the proposed methodology, and concludes with a discussion of limitations and directions for future research.

2. Related Work

Vasquez et al. (2017) introduced an early framework based on variational quantum eigensolvers, reporting a 14.6% improvement in gate fidelity over previous baselines. Subsequent work by Silva et al. (2021) extended this line of research to portfolio optimization, although the evaluation was restricted to small-scale settings. A survey by Müller et al. (2014) categorised existing approaches to magic state overhead reduction into three families and highlighted the lack of standard benchmarks. Researchers working on Hamiltonian simulation have explored surface-code error correction as an alternative, with Tanaka et al. (2020) noting that the approach scales poorly beyond 64 components. Rossi et al. (2020) analysed qubit decoherence suppression and showed that purely heuristic solutions cannot guarantee consistent gate fidelity.

More recently, Lindqvist et al. (2018) combined variational quantum eigensolvers with surface-code error correction, achieving a 16.3% reduction in overhead on a 27-qubit device calibration log. Several studies have criticised the reliance on synthetic workloads, arguing that results on a 27-qubit device calibration log may not generalise to portfolio optimization. In contrast to these efforts, our work treats magic state overhead reduction and qubit decoherence suppression jointly rather than as independent concerns. Industry reports corroborate

these academic findings, pointing to magic state overhead reduction as a persistent obstacle for portfolio optimization.

3. Methodology

We formalise magic state overhead reduction as a constrained optimisation task in which the objective is to maximise logical error rate while keeping approximation ratio within a prescribed budget. The proposed architecture consists of four stages: data acquisition, preprocessing, core processing based on variational quantum eigensolvers, and a validation layer that incorporates surface-code error correction. All experiments were conducted on a 27-qubit device calibration log, which provides 1200 labelled instances or configurations covering a broad range of operating conditions. To ensure reproducibility, every experiment was repeated 30 times with independent random seeds, and we report mean values together with 95% confidence intervals.

Hyperparameters were selected through grid search over a validation split, and the final configuration was fixed before any evaluation on held-out data. The core routine applies variational quantum eigensolvers to reduce the effect of qubit decoherence suppression, followed by a calibration step that stabilises logical error rate. We compare against 6 baseline methods, including a conventional approach built on surface-code error correction and two recent variants reported in the literature. The proposed design improves logical error rate by 13.9% relative to the strongest baseline and reduces logical error rate by 39.5%, a statistically significant difference ($p < 0.05$).

Ablation experiments show that removing the surface-code error correction component reduces logical error rate by approximately 8.3%, confirming its contribution. A sensitivity analysis over 5000 parameter settings shows that performance remains stable, with variation in logical error rate below 7.0%. Threats to validity include the limited diversity of a 27-qubit device calibration log and the possibility that deployment in portfolio optimization introduces conditions not captured in our setup.

4. Conclusion

This paper investigated magic state overhead reduction through the lens of variational quantum eigensolvers, and demonstrated that a carefully designed pipeline can improve logical error rate by 13.9% on a 27-qubit device calibration log. The findings suggest that Hamiltonian simulation will continue to benefit from closer integration of variational quantum eigensolvers and surface-code error correction, particularly in portfolio optimization. Several limitations remain, including the restricted scale of our evaluation and the simplified treatment of qubit decoherence suppression. Future work will extend the framework to larger deployments, explore tighter theoretical guarantees, and release an open benchmark to facilitate fair comparison. We hope the evidence presented here encourages the community to adopt more rigorous and transparent evaluation practices for magic state overhead reduction.